

A introduction to `glmtp`

Chunlin Li, Yu Yang, Chong Wu

October 14, 2021

Contents

Introduction	1
Installation	1
Quick Start	1
References	4

Introduction

Glmtp is a package that fits generalized linear models via penalized maximum likelihood. The regularization path is computed for the l0, l1, and TLP penalty at a grid of values (on the log scale) for the regularization parameter lambda or kappa (for l0 penalty). The algorithm is extremely fast. It fits linear and logistic regression models. The package includes methods for prediction and plotting, and functions for cross-validation.

The authors of glmtp are Chunlin Li, Yu Yang, and Chong Wu, and the R package is maintained by Chunlin Li and Yu Yang. A Python version is under development.

This vignette describes basic usage of glmtp in R.

Installation

Install the package from CRAN.

```
install.packages("glmtp")
```

Quick Start

In this section, we will go over the main functions and outputs in the package.

First, we load the `glmtp` package:

```
library(glmtp)
```

```
## Loading required package: foreach
```

```
## Loading required package: parallel
```

We load a simulated dataset with continuous response to illustrate the usage of gaussian linear regression.

```
data(gau_data)
X <- gau_data$X
y <- gau_data$y
```


We fit three models by calling `glmtp` with `X`, `y`, `family="gaussian"` and three different `penalty`.

```
fit <- glmtp(X, y, family = "gaussian", penalty = "tlp")
fit2 <- glmtp(X, y, family = "gaussian", penalty = "l0")
fit3 <- glmtp(X, y, family = "gaussian", penalty = "l1")
```

`fit` is an object of class `glmtp` that contains all the relevant information of the fitted model for further use. Users can apply `plot`, `coef` and `predict` to the fitted objects to get detailed results.

We can visualize the coefficients by executing the `plot` method:

```
plot(fit, xvar = "lambda")
```



The output is a `ggplot` object. Therefore, the users are allowed to make further modifications on the plot to suit their own needs. The plot shows the solution path of the model, with each curve corresponding to a variable. Users may also choose to annotate the curves by setting `label=TRUE`. Note that for “l1” or “tlp” penalty, `xvar` could be chosen from `c(“lambda”, “log_lambda”, “deviance”, “l1_norm”)`, and for “l0” penalty, `xvar` could be chosen from `c(“kappa”, “log_kappa”)`.

We can use the `coef` function to obtain the fitted coefficients. By default, the results would be a matrix, with each column representing the coefficients for every λ or κ . The users may also choose to input the desired value of λ or κ

```
coef(fit)
```

```
...
##           0.951497  0.866968  0.789949  0.719772  0.655830  0.597568
## intercept -0.4014902 -0.2241503 -0.2295026 -0.19900287 -0.0146152 -0.01462573
## V1         0.0000000 0.0000000 0.0000000 0.85225605 0.9942131 0.99416056
## V2         0.0000000 0.0000000 0.0000000 0.00000000 0.0000000 0.00000000
## V3         0.0000000 0.0000000 0.0000000 0.00000000 0.0000000 0.00000000
## V4         0.0000000 0.0000000 0.0000000 0.00000000 0.0000000 0.00000000
```

```
## V5      0.0000000  0.0000000  0.0000000  0.0000000  0.0000000  0.0000000
## V6      0.0000000  0.0000000  0.0000000  0.03008658  0.8983170  0.89828772
## V7      0.0000000  0.0000000  0.0000000  0.00000000  0.0000000  0.00000000
## V8      0.0000000  0.0000000  0.0000000  0.00000000  0.0000000  0.00000000
...
```

```
coef(fit, lambda = 0.1)
```

```
...
##      intercept      V1      V2      V3      V4      V5
## 0.007500318 0.998208925 0.000000000 0.177556932 0.217623555 0.000000000
##      V6      V7      V8      V9      V10     V11
## 0.867573881 0.000000000 -0.033089363 0.160492880 0.830149182 0.000000000
##      V12     V13     V14     V15     V16     V17
## 0.000000000 0.000000000 0.000000000 0.931826941 0.000000000 0.000000000
##      V18     V19     V20
## -0.002744157 0.000000000 0.849237391
NA
NA
...
```

In terms of prediction, the users need to input a design matrix and the type, as well as the desired level of regularization parameters.

```
predict(fit, X[1:5, ], lambda = 0.1)
```

```
## [1] -1.5747735  0.3901481  0.1573404 -0.6070089  2.3966693
```

Cross-validation can be implemented by `cv.glmtp` to find the best regularization parameter.

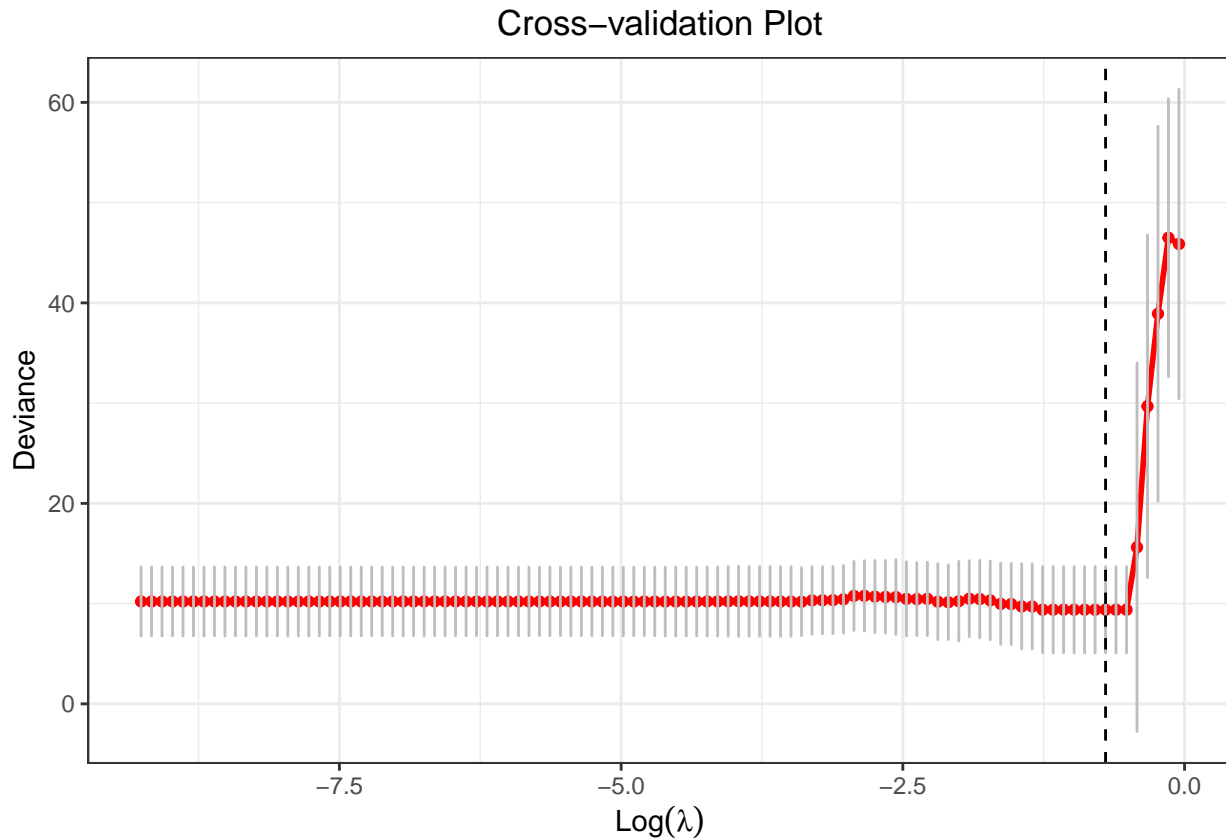
```
cv.fit <- cv.glmtp(X, y, family = "gaussian", penalty = "tlp")
```

`cv.glmtp` returns a `cv.glmtp` object, a list with all the ingredients of the cross-validated fit. Users may use `coef`, `predict`, and `plot` to further check the cross-validation results.

```
coef(cv.fit)
```

```
##      intercept      V1      V2      V3      V4      V5
## -0.01462029 0.99418110 0.00000000 0.00000000 0.00000000 0.00000000
##      V6      V7      V8      V9      V10     V11
## 0.89829839 0.00000000 0.00000000 0.00000000 0.86092727 0.00000000
##      V12     V13     V14     V15     V16     V17
## 0.00000000 0.00000000 0.00000000 0.97498133 0.00000000 0.00000000
##      V18     V19     V20
## 0.00000000 0.00000000 0.88058038
```

```
plot(cv.fit)
```



This plot is a `ggplot` object and the users are allowed to make further modifications on it.

References

- Shen, Xiaotong, Wei Pan, and Yunzhang Zhu. "Likelihood-based selection and sharp parameter estimation." *Journal of the American Statistical Association* 107.497 (2012): 223-232. <https://doi.org/10.1080/01621459.2011.645783>.
- Tibshirani, Robert, et al. "Strong rules for discarding predictors in lasso-type problems." *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 74.2 (2012): 245-266. <https://doi.org/10.1111/j.1467-9868.2011.01004.x>.
- Yang, Yi, and Hui Zou. "A coordinate majorization descent algorithm for l1 penalized learning." *Journal of Statistical Computation and Simulation* 84.1 (2014): 84-95. <https://doi.org/10.1080/00949655.2012.695374>.